

CAP 5636 – Fall 2020 - Homework 2

Step 1: (10 pts)

Solve problems Q1 to Q4 from the Berkeley AI reinforcement learning project:

Q1. Value Iteration

Solution: For this value iteration, we need to implement QValue formula which is as below:

Start with $V_0(s) = 0$: no time steps left means an expected reward sum of zero.

Given vector of $V_k(s)$ values, do one ply of expectimax from each state:

$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Now as to implement this through coding we need to edit below code in **valuelterationAgents.py**:

```
Users > pritijagtap > Downloads > reinforcement 2 > valuelterationAgents.py > ...
40     """
41     self.mdp = mdp
42     self.discount = discount
43     self.iterations = iterations
44     self.values = util.Counter() # A Counter is a dict with default 0
45
46     # Write value iteration code here
47     """ YOUR CODE HERE """
48
49
50     def getValue(self, state):
51         """
52         Return the value of the state (computed in __init__).
53         """
54         return self.values[state]
55
56
57     def computeQValueFromValues(self, state, action):
58         """
59         Compute the Q-value of action in state from the
60         value function stored in self.values.
61         """
62         """ YOUR CODE HERE """
63         util.raiseNotDefined()
64
65
66     def computeActionFromValues(self, state):
67         """
68         The policy is the best action in the given state
69         according to the values currently stored in self.values.
70
71         You may break ties any way you see fit. Note that if
72         there are no legal actions, which is the case at the
73         terminal state, you should return None.
74         """
75         """ YOUR CODE HERE """
76         util.raiseNotDefined()
```

Original code.jpg

After creating the learning code:

```
"""
self.mdp = mdp
self.discount = discount
self.iterations = iterations
self.values = util.Counter() # A Counter is a dict with default 0

# Write value iteration code here
"""*** YOUR CODE HERE ***"""
for _ in range(self.iterations):
    tempvalues = util.Counter()

    for state in self.mdp.getStates():
        if self.mdp.isTerminal(state):
            continue
        actions = self.mdp.getPossibleActions(state)
        value = [self.getQValue(state, action) for action in actions]
        tempvalues[state] = max(value) if value else 0
    self.values = tempvalues
```

```
def computeQValueFromValues(self, state, action):
    """
    Compute the Q-value of action in state from the
    value function stored in self.values.
    """
    """*** YOUR CODE HERE ***"""
    Qvalue = 0
    for next_state_value, transition_prob in self.mdp.getTransitionStatesAndProbs(state, action):
        reward = self.mdp.getReward(state, action, next_state_value)
        future_reward_value = self.discount * self.values[next_state_value]
        Qvalue += transition_prob * (reward + future_reward_value)

    return Qvalue
# util.raiseNotDefined()
```

```

def computeActionFromValues(self, state):
    """
    The policy is the best action in the given state
    according to the values currently stored in self.values.

    You may break ties any way you see fit. Note that if
    there are no legal actions, which is the case at the
    terminal state, you should return None.
    """
    "*** YOUR CODE HERE ***"
    if self.mdp.isTerminal(state):
        return None
    policy = util.Counter()
    for action in self.mdp.getPossibleActions(state):
        policy[action] = self.computeQValueFromValues(state, action)
    return policy.argmax()
    # util.raiseNotDefined()

```

Below is command which we use to run the program:

```
python gridworld.py -a value -i 5
```

After writing code below is the final output:



Output on grid.jpg

Below is Output on terminal:

```

priti jagtap@Pritis-MacBook-Air reinforcement % python gridworld.py -a value -i 5
DEPRECATION WARNING: The system version of Tk is deprecated and may be removed in a future release. Please don't rely on it. Set TK_SILENCE_DEPRECATION=1 to suppress this warning.

RUNNING 1 EPISODES

BEGINNING EPISODE: 1

Started in state: (0, 0)
Took action: north
Ended in state: (0, 1)
Got reward: 0.0

Started in state: (0, 1)
Took action: north
Ended in state: (0, 2)
Got reward: 0.0

Started in state: (0, 2)
Took action: east
Ended in state: (1, 2)
Got reward: 0.0

Started in state: (1, 2)
Took action: east
Ended in state: (2, 2)
Got reward: 0.0

Started in state: (2, 2)
Took action: east
Ended in state: (3, 2)
Got reward: 0.0

Started in state: (3, 2)
Took action: exit
Ended in state: TERMINAL_STATE
Got reward: 1

EPISODE 1 COMPLETE: RETURN WAS 0.59049

AVERAGE RETURNS FROM START STATE: 0.59049

```

If we run the autograder command get the below output:

```

priti jagtap@Pritis-MacBook-Air reinforcement % python autograder.py -q q1
Starting on 11-13 at 23:33:08

Question q1
=====

*** PASS: test_cases/q1/1-tinygrid.test
*** PASS: test_cases/q1/2-tinygrid-noisy.test
*** PASS: test_cases/q1/3-bridge.test
*** PASS: test_cases/q1/4-discountgrid.test

### Question q1: 6/6 ###

Finished at 23:33:08

Provisional grades
=====
Question q1: 6/6
-----

Total: 6/6

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.

```

For this Question we got below values against each iteration and the graph is plotted in Step 2.

Iteration	Values
1	0
2	0
3	0
4	0.37
5	0.5
6	0.58
7	0.61
8	0.63
9	0.64
10	0.64

Q2. Bridge Crossing Analysis

For this question we have to change analysis.py. Change the values for answerDiscount = 0.9 and answerNoise = 0.2.

Original code:

```
#####  
# ANALYSIS QUESTIONS #  
#####  
  
# Set the given parameters to obtain the specified policies through  
# value iteration.  
  
def question2():  
    answerDiscount = 0.9  
    answerNoise = 0.2  
    return answerDiscount, answerNoise
```

We tried lot of combinations for answerDiscount and answerNoise values. Like answerDiscount = 0.1 and answerNoise = 0.5. For mentioned values below is output and we used this command:
`python gridworld.py -a value -i 100 -g BridgeGrid --discount 0.1 --noise 0.5`

```
def question2():
    answerDiscount = 0.1
    answerNoise = 0.5
    return answerDiscount, answerNoise
```



After creating the learning code:

```
#####
# ANALYSIS QUESTIONS #
#####

# Set the given parameters to obtain the specified policies through
# value iteration.

def question2():
    answerDiscount = 0.9
    answerNoise = 0.0
    return answerDiscount, answerNoise
```

We are using below command to get the output:

```
python gridworld.py -a value -i 100 -g BridgeGrid --discount 0.9 --noise 0.0
```

This is the final output for the program:



Q3: Policies

In this question we need to use the optimal policy types for below. We are editing **analysis.py**, in the file we going to edit **answerDiscount**, **answerNoise** and **answerLivingReward** paramters in the file.

- Prefer the close exit (+1), risking the cliff (-10)
For this we are editing def question3a:

Original code:

```
def question3a():  
    answerDiscount = None  
    answerNoise = None  
    answerLivingReward = None  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```


After creating the learning code:

```
def question3a():
    answerDiscount = 0.2
    answerNoise = 0.2
    answerLivingReward = -7
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'
```

- b. Prefer the close exit (+1), but avoiding the cliff (-10)
For this we are editing def question3b:

Original code:

```
def question3b():
    answerDiscount = None
    answerNoise = None
    answerLivingReward = None
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'
```

After creating the learning code:

```
def question3b():
    answerDiscount = 0.2
    answerNoise = 0.2
    answerLivingReward = -0.4
    return answerDiscount, answerNoise, answerLivingReward
    # If not possible, return 'NOT POSSIBLE'
```

- c. Prefer the distant exit (+10), risking the cliff (-10)
For this we are editing def question3c:

Original code:

```
def question3c():  
    answerDiscount = None  
    answerNoise = None  
    answerLivingReward = None  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```

After creating the learning code:

```
def question3c():  
    answerDiscount = 0.9  
    answerNoise = 0.0  
    answerLivingReward = 0.0  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```

- d. Prefer the distant exit (+10), avoiding the cliff (-10)

For this we are editing def question3d:

Original code:

```
def question3d():  
    answerDiscount = None  
    answerNoise = None  
    answerLivingReward = None  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```

After creating the learning code:

```
def question3d():  
    answerDiscount = 0.8  
    answerNoise = 0.3  
    answerLivingReward = 0.3  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```

- e. Avoid both exits and the cliff (so an episode should never terminate)
For this we are editing def question3e:

Original code:

```
def question3e():  
    answerDiscount = None  
    answerNoise = None  
    answerLivingReward = None  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```

After creating the learning code:

```
def question3e():  
    answerDiscount = 0.9  
    answerNoise = 0.0  
    answerLivingReward = 1.0  
    return answerDiscount, answerNoise, answerLivingReward  
    # If not possible, return 'NOT POSSIBLE'
```

Once done editing the analysis.py we are going to run code using below command:

```
python autograder.py -q q3
```

Below is the output for the program:

```

pritijagtap@Pritis-MacBook-Air reinforcement % python autograder.py -q q3
Starting on 11-15 at 22:00:53

Question q3
=====

*** PASS: test_cases/q3/1-question-3.1.test
*** PASS: test_cases/q3/2-question-3.2.test
*** PASS: test_cases/q3/3-question-3.3.test
*** PASS: test_cases/q3/4-question-3.4.test
*** PASS: test_cases/q3/5-question-3.5.test

### Question q3: 5/5 ###

Finished at 22:00:54

Provisional grades
=====
Question q3: 5/5
-----
Total: 5/5

Your grades are NOT yet registered. To register your grades, make sure
to follow your instructor's guidelines to receive credit on your project.

```

Q4: Q-Learning

For this question we have edit 4 functions as mentioned in the instructions from file qlearningAgents.py and functions are update, computeValueFromQValues, getQValue, and computeActionFromQValues:

Original code:

```

def __init__(self, **args):
    "You can initialize Q-values here..."
    ReinforcementAgent.__init__(self, **args)

    "*** YOUR CODE HERE ***"

```

```
def getQValue(self, state, action):  
    """  
    Returns Q(state,action)  
    Should return 0.0 if we have never seen a state  
    or the Q node value otherwise  
    """  
    "*** YOUR CODE HERE ***"  
    util.raiseNotDefined()
```

```
def computeValueFromQValues(self, state):  
    """  
    Returns max_action Q(state,action)  
    where the max is over legal actions. Note that if  
    there are no legal actions, which is the case at the  
    terminal state, you should return a value of 0.0.  
    """  
    "*** YOUR CODE HERE ***"  
    util.raiseNotDefined()
```

```
def computeActionFromQValues(self, state):  
    """  
    Compute the best action to take in a state. Note that if there  
    are no legal actions, which is the case at the terminal state,  
    you should return None.  
    """  
    "*** YOUR CODE HERE ***"  
    util.raiseNotDefined()
```

```
def update(self, state, action, nextState, reward):
    """
    The parent class calls this to observe a
    state = action => nextState and reward transition.
    You should do your Q-Value update here

    NOTE: You should never call this function,
    it will be called on your behalf
    """
    "*** YOUR CODE HERE ***"
    util.raiseNotDefined()
```

After creating the learning code:

```
def __init__(self, **args):
    "You can initialize Q-values here..."
    ReinforcementAgent.__init__(self, **args)

    "*** YOUR CODE HERE ***"
    self.qvalue = util.Counter()
```

```
def getQValue(self, state, action):
    """
    Returns Q(state,action)
    Should return 0.0 if we have never seen a state
    or the Q node value otherwise
    """
    "*** YOUR CODE HERE ***"
    return self.qvalue[(state, action)]
    # util.raiseNotDefined()
```

```

def computeValueFromQValues(self, state):
    """
    Returns max_action Q(state,action)
    where the max is over legal actions. Note that if
    there are no legal actions, which is the case at the
    terminal state, you should return a value of 0.0.
    """
    "*** YOUR CODE HERE ***"
    if len(self.getLegalActions(state)) == 0:
        return 0.0
    maximum_qvalue = float('-inf')
    for actions in self.getLegalActions(state):
        maximum_qvalue = max(maximum_qvalue, self.getQValue(state, actions))
    return maximum_qvalue
# util.raiseNotDefined()

```

```

def computeActionFromQValues(self, state):
    """
    Compute the best action to take in a state. Note that if there
    are no legal actions, which is the case at the terminal state,
    you should return None.
    """
    "*** YOUR CODE HERE ***"
    if len(self.getLegalActions(state)) == 0:
        return None
    bestaction_in_state = []
    best_qvalue = self.computeValueFromQValues(state)
    for actions in self.getLegalActions(state):
        if best_qvalue == self.getQValue(state, actions):
            bestaction_in_state.append(actions)

    return random.choice(bestaction_in_state)
# util.raiseNotDefined()

```

```

def update(self, state, action, nextState, reward):
    """
    The parent class calls this to observe a
    state = action => nextState and reward transition.
    You should do your Q-Value update here

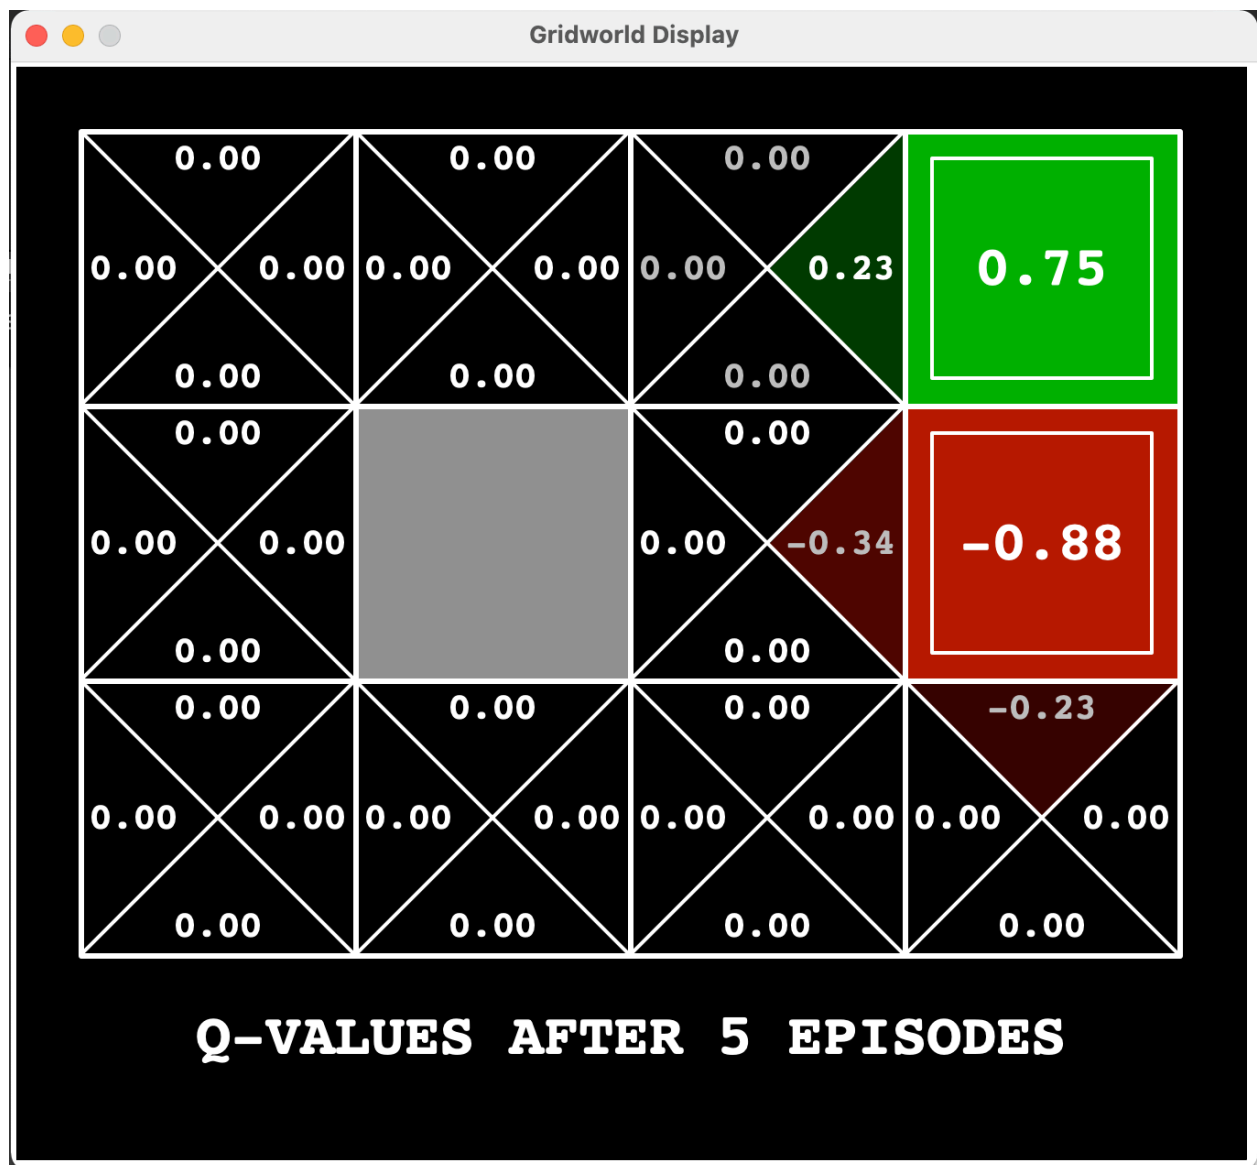
    NOTE: You should never call this function,
    it will be called on your behalf
    """
    "*** YOUR CODE HERE ***"
    update_q_value = self.qvalue[(state, action)]
    values = reward + (self.discount * self.computeValueFromQValues(nextState))
    self.qvalue[(state, action)] = (1 - self.alpha) * update_q_value + self.alpha * values
    # util.raiseNotDefined()

```

Final output on gridworld:


After running program for 100 iterations, we got below output:

We got different values through the 5 iterations for each episode and those values are stored in one csv file.



Iteration	North	East	West	South	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	
5	0	0	0	0	
6	0	0	0	0	
7	0	0	0	0	
8	0	0	0	0	
9	0	0	0	0.056429	
10	0	0	0	0.056429	
11	0	0	0	0.056429	
12	0	0	0	0.056429	
13	0	0	0	0.056429	
14	0	0	0	0.056429	
15	0	0	0	0.056429	
16	0	0	0	0.056429	
17	0	0	0	0.056429	
18	0	0	0	0.056429	
19	0	0	0	0.056429	
20	0	0	0	0.056429	
21	0	0	0	0.056429	
22	0	0	0	0.056429	
23	0	0	0	0.056429	
24	0	0	0	0.056429	

Step 2: Visualize the evolution of the V-value (4 pts)



Step 3: Visualize the evolution of the Q-value (4 pts)

